

PERFORMANCE EVALUATION OF OPENMPI COLLECTIVE OPERATIONS

**Pranay Narsipuram
SM3800006**

JULY 2024

OBJECTIVE

1

To evaluate and compare the performance of various algorithms in the OpenMPI library for collective operations, specifically focusing on broadcast and gather operations.

2

To develop and validate performance models that predict the latency of these implementations.

EXPERIMENTAL SETUP

- ORFEO cluster: 2 EPYC nodes, 128 cores per node.
- OpenMPI 4.1.5
- OSU Micro-Benchmarks (osu_bcast, osu_gather)
- Scripts to automate data collection.
- –map-by core policy.
- Collect measures varying the number of processes from 2 to 256 and the size of the messages from 1 to 524288 bytes.




DATA COLLECTION




scan for
benchmark_job.sh

- Usage of benchmark_job.sh script.
- Usage: `./benchmark_job.sh <operation> <benchmark_type>`
- Operations: Broadcast and Gather.
- Algorithms: ignore, chain, pipeline, binary_tree, basic_linear, binomial, linear_sync.
- Fixed and variable message sizes.

MPI_Bcast



MPI_Gather



OSU MICRO BENCHMARKS

- A suite of benchmarks developed by Ohio State University.
- Key Benchmarks:
 - `osu_latency`: Measures point-to-point latency.
 - `osu_bw`: Measures point-to-point bandwidth.
 - `osu_bcast`: Measures performance of MPI_Bcast.
 - `osu_gather`: Measures performance of MPI_Gather.
- Benchmarking Methodology:
 - Varying number of processes and message sizes.
- Automated benchmarking scripts.


Broadcast Operation

Evaluated the performance of different broadcast algorithms using OpenMPI on the ORFEO cluster.


- **Ignore Algorithm:**
Baseline for comparison, essentially a no-operation broadcast.
Results: Minimal latency, serves as a control to compare other algorithms.
- **Chain Algorithm:**
Each process forwards the message to the next sequentially.
Results: Higher latency due to sequential forwarding. Suitable for small-scale systems.
Analysis: Latency increases linearly with the number of processes.
- **Pipeline Algorithm:**
Processes arranged in a pipeline, reducing latency by overlapping communication with computation.
Results: Better performance than the chain algorithm due to overlapping.
Analysis: Latency improves but increases with larger message sizes.
- **Binary Tree Algorithm:**
Root sends messages to two children, forming a binary tree.
Results: Best performance among tested algorithms. Efficient for larger process counts due to logarithmic scaling.
Analysis: Latency increases logarithmically with the number of processes, making it highly scalable.
- Comparison:
- **Fixed Message Size:** Analyzed latency with a fixed message size across different algorithms.
Results: Binary tree algorithm consistently outperforms others.
- **Variable Message Size:** Analyzed latency changes with varying message sizes.
Results: Pipeline and binary tree algorithms show better scalability with increasing message sizes.

Visualizing the performance of broadcast algorithms using 3D heatmaps on the ORFEO cluster.


Broadcast Latency
map-by core, ignore algorithm




Broadcast Latency
map-by core, chain algorithm



Broadcast Latency
map-by core, pipeline algorithm



Broadcast Latency
map-by core, binary_tree algorithm




Gather Operation

Evaluated the performance of different gather algorithms using OpenMPI on the ORFEO cluster.


- **Ignore Algorithm:**
Baseline for comparison, essentially a no-operation gather.
Results: Minimal latency, serves as a control to compare other algorithms.
- **Basic Linear Algorithm:**
Each process sends its data directly to the root.
Results: Simple implementation. High latency due to all processes sending data to the root.
Analysis: Latency increases linearly with the number of processes.
- **Binomial Algorithm:**
Data is aggregated in a hierarchical binomial tree structure.
Results: Better performance than the basic linear algorithm due to hierarchical aggregation.
Analysis: Latency scales logarithmically with the number of processes.
- **Linear Sync Algorithm:**
Similar to basic linear but includes synchronization steps.
Results: Ensures data consistency and reduces network congestion.
Analysis: Slightly higher overhead due to synchronization, but better performance in terms of consistency.
- **Comparison:**
 - Fixed Message Size:** Analyzed latency with a fixed message size across different algorithms.
Results: Binomial algorithm consistently outperforms others in fixed message size scenarios.
 - Variable Message Size:** Analyzed latency changes with varying message sizes.
Results: Binomial and linear sync algorithms show better scalability with increasing message sizes.

Visualizing the performance of gather algorithms using 3D heatmaps on the ORFEO cluster.


Gather Latency
map-by core, ignore algorithm




Gather Latency
map-by core, binomial algorithm



Gather Latency
map-by core, basic_linear algorithm



Gather Latency
map-by core, linear_sync algorithm




Performance Model - Broadcast fixed

- Model shows strong predictive power, closely matching actual latencies.
- Effective across different core counts, ensuring reliability.

Model Equation:

$$\log_2(\text{avg-lat}) = \beta_1 \cdot \text{proc-num} + \beta_2 \cdot \log_2(\text{mess-size}) + \beta_3 \cdot (\log_2(\text{mess-size}))^2$$

- Polynomial term captures non-linear trends in latency.



Performance Model - Gather Fixed

- Model shows strong predictive power, closely matching actual latencies.
- Effective across different core counts, ensuring reliability.

Model Equation:

$$\log_2(\text{avg-lat}) = \delta_1 \cdot \text{proc-num} + \delta_2 \cdot \log_2(\text{mess-size}) + \delta_3 \cdot (\log_2(\text{mess-size}))^2$$

- Polynomial term captures non-linear trends in latency.

